

# Natural Language Processing (CSE 517 & 447): Multinomial Logistic Regression

Noah Smith

© 2022

University of Washington  
nasmith@cs.washington.edu

Winter 2022

Readings: Eisenstein (2019) 2 and Appendix B

# Motivation

- ▶ Dominant perspective in NLP in the 1990s–today: supervised machine learning
  - ▶ This lecture's model is a direct ancestor of today's popular methods.
- ▶ Engineering approach: feature design
- ▶ Relevance today: interpretable and efficient classification

# Classification in NLP

We approach many problems in NLP by treating them as problems of classification.

- ▶ Input might be a document, a paragraph, a sentence, a word
- ▶ Output is a label from a finite set of classes or labels, denoted  $\mathcal{L}$ , defined by your application or theory

Notation:  $\text{classify} : \mathcal{V}^* \rightarrow \mathcal{L}$  is a classifier, e.g., the one you build. It is deterministic and typically constructed from data and machine learning.

# Text (Document) Classification Examples

- ▶ Library-like subjects (e.g., the Dewey decimal system)
- ▶ News stories: politics vs. sports vs. business vs. technology ...
- ▶ Reviews of films, restaurants, products: positive vs. negative
- ▶ Author attributes: identity, political stance, gender, age, ...
- ▶ Email, arXiv submissions, etc.: spam vs. not
- ▶ What is the reading level of a piece of text?
- ▶ How influential will a scientific paper be?
- ▶ Will a piece of proposed legislation pass?
- ▶ What dialect is a text written in?
- ▶ Does the text contain content that will likely offend people?

# Notation

$\mathcal{V}$  is the set of words in the language we're working with.

$\mathbf{X}$  is a random variable for texts (inputs); in a given instance it takes a value from  $\mathcal{V}^*$  (sequences of words).

$Y$  is a random variable for labels (outputs); in a given instance it takes a value from  $\mathcal{L}$ .

$p(\mathbf{X}, Y)$  is the “true” distribution of labeled texts;  $p(Y)$  is the distribution of labels. **Normally, we do not know these distributions except by looking at data.**

# Evaluating a Classifier

Accuracy:

$$\begin{aligned} A(\text{classify}) &= p(\text{classify}(\mathbf{X}) = Y) \\ &= \sum_{\mathbf{x} \in \mathcal{V}^*, \ell \in \mathcal{L}} p(\mathbf{X} = \mathbf{x}, Y = \ell) \cdot \begin{cases} 1 & \text{if } \text{classify}(\mathbf{x}) = \ell \\ 0 & \text{otherwise} \end{cases} \\ &= \sum_{\mathbf{x} \in \mathcal{V}^*, \ell \in \mathcal{L}} p(\mathbf{X} = \mathbf{x}, Y = \ell) \cdot \mathbf{1}\{\text{classify}(\mathbf{x}) = \ell\} \end{aligned}$$

where  $p$  is the *true* distribution over data. Error is  $1 - A$ .

# Evaluating a Classifier

Accuracy:

$$\begin{aligned} A(\text{classify}) &= p(\text{classify}(\mathbf{X}) = Y) \\ &= \sum_{\mathbf{x} \in \mathcal{V}^*, \ell \in \mathcal{L}} p(\mathbf{X} = \mathbf{x}, Y = \ell) \cdot \begin{cases} 1 & \text{if } \text{classify}(\mathbf{x}) = \ell \\ 0 & \text{otherwise} \end{cases} \\ &= \sum_{\mathbf{x} \in \mathcal{V}^*, \ell \in \mathcal{L}} p(\mathbf{X} = \mathbf{x}, Y = \ell) \cdot \mathbf{1} \{\text{classify}(\mathbf{x}) = \ell\} \end{aligned}$$

where  $p$  is the *true* distribution over data. Error is  $1 - A$ .

This is *estimated* using a test dataset  $\langle \bar{\mathbf{x}}_1, \bar{y}_1 \rangle, \dots, \langle \bar{\mathbf{x}}_m, \bar{y}_m \rangle$ :

$$\hat{A}(\text{classify}) = \frac{1}{m} \sum_{i=1}^m \mathbf{1} \{\text{classify}(\bar{\mathbf{x}}_i) = \bar{y}_i\}$$

# Some Issues with Test-Set Accuracy



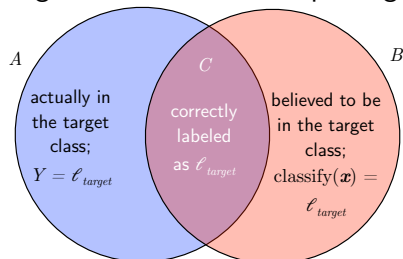
# Some Issues with Test-Set Accuracy

- ▶ Class imbalance: if  $p(Y = \text{not spam}) = 0.99$ , then you can get  $\hat{A} \approx 0.99$  by always guessing “not spam.”

# Evaluation in the “Needle in a Haystack” Case

Suppose one label  $\ell_{target} \in \mathcal{L}$  is a “target.”

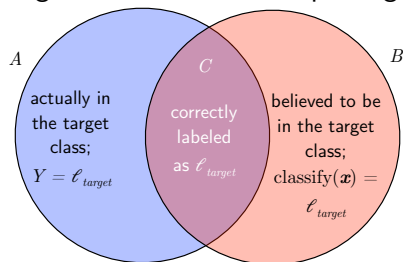
**Precision** and **recall** encode the goals of returning a “pure” set of targeted instances and capturing *all* of them.



# Evaluation in the “Needle in a Haystack” Case

Suppose one label  $\ell_{target} \in \mathcal{L}$  is a “target.”

**Precision** and **recall** encode the goals of returning a “pure” set of targeted instances and capturing *all* of them.

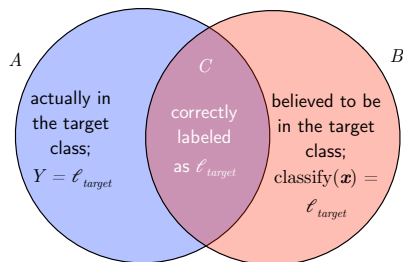


$$\hat{P}(\text{classify}) = \frac{|C|}{|B|} = \frac{|A \cap B|}{|B|}$$

$$\hat{R}(\text{classify}) = \frac{|C|}{|A|} = \frac{|A \cap B|}{|A|}$$

$$\hat{F}_1(\text{classify}) = 2 \cdot \frac{\hat{P} \cdot \hat{R}}{\hat{P} + \hat{R}}$$

## Another View: Contingency Table



	$Y = \ell_{target}$	$Y \neq \ell_{target}$	
$\text{classify}(\mathbf{X}) = \ell_{target}$	$ C $ (true positives)	$ B \setminus C $ (false positives)	$ B $
$\text{classify}(\mathbf{X}) \neq \ell_{target}$	$ A \setminus C $ (false negatives)	(true negatives)	
	$ A $		

# Generalization of Precision and Recall

Macroaveraged precision and recall: let each class be the “target” and report the average  $\hat{P}$  and  $\hat{R}$  across all classes.

Microaveraged precision and recall: pool all one-vs.-rest decisions into a single contingency table, calculate  $\hat{P}$  and  $\hat{R}$  from that.

# Some Issues with Test-Set Accuracy

- ▶ Class imbalance: if  $p(Y = \text{not spam}) = 0.99$ , then you can get  $\hat{A} \approx 0.99$  by always guessing “not spam.”
  - ▶ Solution: report precision and recall

# Some Issues with Test-Set Accuracy

- ▶ Class imbalance: if  $p(Y = \text{not spam}) = 0.99$ , then you can get  $\hat{A} \approx 0.99$  by always guessing “not spam.”
  - ▶ Solution: report precision and recall
- ▶ Relative importance of classes or cost of error types.

# Some Issues with Test-Set Accuracy

- ▶ Class imbalance: if  $p(Y = \text{not spam}) = 0.99$ , then you can get  $\hat{A} \approx 0.99$  by always guessing “not spam.”
  - ▶ Solution: report precision and recall
- ▶ Relative importance of classes or cost of error types.
  - ▶ Solution: report precision and recall for each class, or categorize different error types.



# Some Issues with Test-Set Accuracy

- ▶ Class imbalance: if  $p(Y = \text{not spam}) = 0.99$ , then you can get  $\hat{A} \approx 0.99$  by always guessing “not spam.”
  - ▶ Solution: report precision and recall
- ▶ Relative importance of classes or cost of error types.
  - ▶ Solution: report precision and recall for each class, or categorize different error types.
- ▶ Variance due to the test data.

# Some Issues with Test-Set Accuracy

- ▶ Class imbalance: if  $p(Y = \text{not spam}) = 0.99$ , then you can get  $\hat{A} \approx 0.99$  by always guessing “not spam.”
  - ▶ Solution: report precision and recall
- ▶ Relative importance of classes or cost of error types.
  - ▶ Solution: report precision and recall for each class, or categorize different error types.
- ▶ Variance due to the test data.
  - ▶ Solution: repeat entire experiment with shuffled data, multiple times, and report mean and standard deviation.

# Some Issues with Test-Set Accuracy

- ▶ Class imbalance: if  $p(Y = \text{not spam}) = 0.99$ , then you can get  $\hat{A} \approx 0.99$  by always guessing “not spam.”
  - ▶ Solution: report precision and recall
- ▶ Relative importance of classes or cost of error types.
  - ▶ Solution: report precision and recall for each class, or categorize different error types.
- ▶ Variance due to the test data.
  - ▶ Solution: repeat entire experiment with shuffled data, multiple times, and report mean and standard deviation.
  - ▶ Test data is not representative of real data.

(Some additional topics in the “extras” section at the end of this file: cross-validation and statistical significance.)

# Building a Text Classifier: Standard Line of Attack

1. Human experts label some data, or nature provides labeled data.
2. Feed the data to a supervised machine learning algorithm that constructs an automatic classifier  $\text{classify} : \mathcal{V}^* \rightarrow \mathcal{L}$
3. Apply  $\text{classify}$  to as much data as you want!

Note: we assume the texts are segmented into symbols from  $\mathcal{V}$ , even the new ones.

# Features of a Text

Running example:

$x =$  "The vodka was great, but don't touch the hamburgers."

A different representation of the text sequences: features.

# Features of a Text

Running example:

$\mathbf{x}$  = "The vodka was great, but don't touch the hamburgers."

A different representation of the text sequences: features.

- Often, these are term (word or word sequence) frequencies.

E.g.,  $\phi_{\text{hamburgers}}^{\text{freq.}}(\mathbf{x}) = 1$ ,  $\phi_{\text{the}}^{\text{freq.}}(\mathbf{x}) = 2$ ,  $\phi_{\text{delicious}}^{\text{freq.}}(\mathbf{x}) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{freq.}}(\mathbf{x}) = 1$ .

# Features of a Text

Running example:

$x = \text{"The vodka was great, but don't touch the hamburgers."}$

A different representation of the text sequences: features.

- Often, these are term (word or word sequence) frequencies.

E.g.,  $\phi_{\text{hamburgers}}^{\text{freq.}}(x) = 1$ ,  $\phi_{\text{the}}^{\text{freq.}}(x) = 2$ ,  $\phi_{\text{delicious}}^{\text{freq.}}(x) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{freq.}}(x) = 1$ .

“Bag of words” model: one based on word frequency features alone.





# Features of a Text

Running example:

$\mathbf{x}$  = "The vodka was great, but don't touch the hamburgers."

A different representation of the text sequences: features.

- Often, these are term (word or word sequence) frequencies.

E.g.,  $\phi_{\text{hamburgers}}^{\text{freq.}}(\mathbf{x}) = 1$ ,  $\phi_{\text{the}}^{\text{freq.}}(\mathbf{x}) = 2$ ,  $\phi_{\text{delicious}}^{\text{freq.}}(\mathbf{x}) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{freq.}}(\mathbf{x}) = 1$ .

# Features of a Text

Running example:

$x$  = "The vodka was great, but don't touch the hamburgers."

A different representation of the text sequences: features.

- Often, these are term (word or word sequence) frequencies.

E.g.,  $\phi_{\text{hamburgers}}^{\text{freq.}}(x) = 1$ ,  $\phi_{\text{the}}^{\text{freq.}}(x) = 2$ ,  $\phi_{\text{delicious}}^{\text{freq.}}(x) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{freq.}}(x) = 1$ .

- Can also be binary word "presence" features.

E.g.,  $\phi_{\text{hamburgers}}^{\text{presence}}(x) = 1$ ,  $\phi_{\text{the}}^{\text{presence}}(x) = 1$ ,  $\phi_{\text{delicious}}^{\text{presence}}(x) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{presence}}(x) = 1$ .

# Features of a Text

Running example:

$\mathbf{x}$  = "The vodka was great, but don't touch the hamburgers."

A different representation of the text sequences: features.

- ▶ Often, these are term (word or word sequence) frequencies.  
E.g.,  $\phi_{\text{hamburgers}}^{\text{freq.}}(\mathbf{x}) = 1$ ,  $\phi_{\text{the}}^{\text{freq.}}(\mathbf{x}) = 2$ ,  $\phi_{\text{delicious}}^{\text{freq.}}(\mathbf{x}) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{freq.}}(\mathbf{x}) = 1$ .
- ▶ Can also be binary word "presence" features.  
E.g.,  $\phi_{\text{hamburgers}}^{\text{presence}}(\mathbf{x}) = 1$ ,  $\phi_{\text{the}}^{\text{presence}}(\mathbf{x}) = 1$ ,  $\phi_{\text{delicious}}^{\text{presence}}(\mathbf{x}) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{presence}}(\mathbf{x}) = 1$ .
- ▶ Transformations on word frequencies: logarithm, idf weighting

$$\forall v \in \mathcal{V}, \text{idf}(v) = \log \frac{n}{|i : \text{count}_{\mathbf{x}_i}(v) > 0|}$$

$$\phi_v^{\text{tfidf}}(\mathbf{x}) = \phi_v^{\text{freq.}}(\mathbf{x}) \cdot \text{idf}(v)$$

# Features of a Text

Running example:

$\mathbf{x}$  = "The vodka was great, but don't touch the hamburgers."

A different representation of the text sequences: features.

- Often, these are term (word or word sequence) frequencies.

E.g.,  $\phi_{\text{hamburgers}}^{\text{freq.}}(\mathbf{x}) = 1$ ,  $\phi_{\text{the}}^{\text{freq.}}(\mathbf{x}) = 2$ ,  $\phi_{\text{delicious}}^{\text{freq.}}(\mathbf{x}) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{freq.}}(\mathbf{x}) = 1$ .

- Can also be binary word "presence" features.

E.g.,  $\phi_{\text{hamburgers}}^{\text{presence}}(\mathbf{x}) = 1$ ,  $\phi_{\text{the}}^{\text{presence}}(\mathbf{x}) = 1$ ,  $\phi_{\text{delicious}}^{\text{presence}}(\mathbf{x}) = 0$ ,  
 $\phi_{\text{don't touch}}^{\text{presence}}(\mathbf{x}) = 1$ .

- Transformations on word frequencies: logarithm, idf weighting

$$\forall v \in \mathcal{V}, \text{idf}(v) = \log \frac{n}{|i : \text{count}_{\mathbf{x}_i}(v) > 0|}$$

$$\phi_v^{\text{tfidf}}(\mathbf{x}) = \phi_v^{\text{freq.}}(\mathbf{x}) \cdot \text{idf}(v)$$

- "Bias" feature,  $\phi^{\text{bias}}$  which takes a constant value of 1.

# Reflection

Given what you already know about words, can you think of features that might generalize better than the ones just discussed?

# Features are Extremely Important!

The features fully determine what a learned model “sees” about an example.

We often stack the features into a **feature vector**:  $\phi(\mathbf{x}) \in \mathbb{R}^d$ , which “embeds” the input  $\mathbf{x}$  in  $d$ -dimensional space

# Aperitif: (Binary) Logistic Regression

A logistic regression model is defined by:

- ▶ A collection of feature functions, denoted  $\phi_1, \dots, \phi_d$ , each mapping  $\mathcal{V}^* \rightarrow \mathbb{R}$ .
  - ▶ **The designer of the system chooses the features.**
- ▶ A coefficient or “weight” for every feature, denoted  $\theta_1, \dots, \theta_d$ , each  $\in \mathbb{R}$ .
  - ▶ **The weights are “parameters” that are chosen automatically by applying a learning algorithm.**


The label set is  $\mathcal{L} = \{+1, -1\}$ .

$$\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}) = \sum_{j=1}^d \theta_j \phi_j(\mathbf{x}) = \boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x})$$


$$\text{classify}_{\text{LR}}(\mathbf{x}) = \text{sign}(\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}))$$



# Computation Graph View of LR Classifier



# Geometric View of LR



# Learning a Logistic Regression Classifier


Learning requires us to choose the weight vector,  $\theta$ .

There are many ways you could do this; logistic regression tells you what vector you should choose based on a probabilistic view of the classifier (but not exactly *how* to find it).

# Reflection

Recall the bias feature,  $\phi^{bias}(\mathbf{x}) = 1$ . What role does it play in the geometric interpretation of the model?

# Standard Logistic Function



$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

# Probabilistic View of LR

Our model actually defines a *probability* distribution over the labels  $\mathcal{L} = \{+1, -1\}$ :

$$p_{\text{LR}}(Y = +1 \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \sigma(\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}))$$

# Probabilistic View of LR

Our model actually defines a *probability* distribution over the labels  $\mathcal{L} = \{+1, -1\}$ :

$$p_{\text{LR}}(Y = +1 \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \sigma(\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}))$$

$$p_{\text{LR}}(Y = -1 \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = 1 - \sigma(\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta})) = \sigma(-\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}))$$

# Probabilistic View of LR

Our model actually defines a *probability* distribution over the labels  $\mathcal{L} = \{+1, -1\}$ :

$$p_{\text{LR}}(Y = +1 \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \sigma(\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}))$$

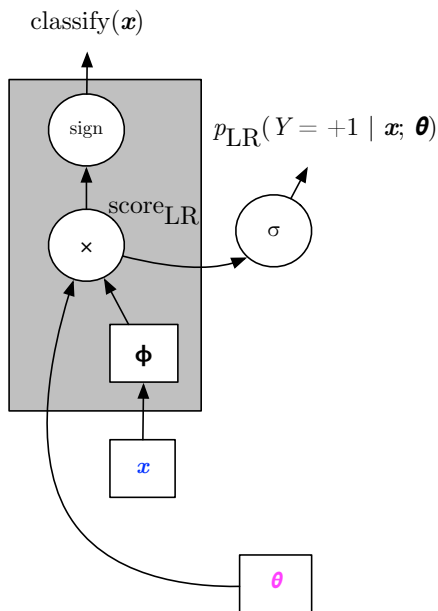
$$p_{\text{LR}}(Y = -1 \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = 1 - \sigma(\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta})) = \sigma(-\text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}))$$

$$p_{\text{LR}}(Y = y \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \sigma(y \cdot \text{score}_{\text{LR}}(\mathbf{x}; \boldsymbol{\theta}))$$

Note: recorded lecture has a mistake on the line above (at 47:35); there should *not* be a minus sign in front of  $y$ .



# Computation Graph View of LR Probability



# Probabilistic View of LR

This suggests using the principle of maximum likelihood to estimate  $\theta$ :

$$\theta^* = \arg \max_{\theta \in \mathbb{R}^d} \prod_{i=1}^n p_{\text{LR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta)$$

# Probabilistic View of LR

This suggests using the principle of maximum likelihood to estimate  $\theta$ :


$$\begin{aligned}\theta^* &= \arg \max_{\theta \in \mathbb{R}^d} \prod_{i=1}^n p_{\text{LR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta) \\ &= \arg \max_{\theta \in \mathbb{R}^d} \sum_{i=1}^n \log p_{\text{LR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta)\end{aligned}$$

# Probabilistic View of LR

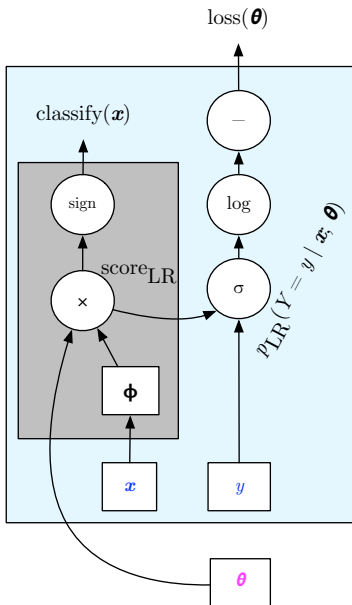
This suggests using the principle of maximum likelihood to estimate  $\theta$ :

$$\begin{aligned}\theta^* &= \arg \max_{\theta \in \mathbb{R}^d} \prod_{i=1}^n p_{\text{LR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta) \\ &= \arg \max_{\theta \in \mathbb{R}^d} \sum_{i=1}^n \log p_{\text{LR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta) \\ &= \arg \min_{\theta \in \mathbb{R}^d} \sum_{i=1}^n \underbrace{-\log p_{\text{LR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta)}_{\text{sometimes called "log loss" or "cross entropy"}}\end{aligned}$$


# Computation Graph View of LR Probability of Correct Label $y$



# Computation Graph View of Log Loss (One Instance)



# Computation Graph View of Log Loss (Many Instances)



# Learning for Logistic Regression

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta} \in \mathbb{R}^d} \underbrace{\sum_{i=1}^n \log \left( 1 + \exp \left( -y_i \cdot \boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x}_i) \right) \right)}_{\text{loss}(\boldsymbol{\theta})}$$



# Learning for Logistic Regression

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta} \in \mathbb{R}^d} \underbrace{\sum_{i=1}^n \log \left( 1 + \exp \left( -y_i \cdot \boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x}_i) \right) \right)}_{\text{loss}(\boldsymbol{\theta})}$$

- You can efficiently implement the objective function “loss” given your data and your features  $\boldsymbol{\phi}$ .

# Learning for Logistic Regression

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta} \in \mathbb{R}^d} \underbrace{\sum_{i=1}^n \log \left( 1 + \exp \left( -y_i \cdot \boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x}_i) \right) \right)}_{\text{loss}(\boldsymbol{\theta})}$$

- ▶ You can efficiently implement the objective function “loss” given your data and your features  $\boldsymbol{\phi}$ .
- ▶ Because it is continuous and differentiable, and the optimization problem is unconstrained, you can use the *gradient* of loss to iteratively move closer to a minimum.

# Learning for Logistic Regression

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta} \in \mathbb{R}^d} \underbrace{\sum_{i=1}^n \log \left( 1 + \exp \left( -y_i \cdot \boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x}_i) \right) \right)}_{\text{loss}(\boldsymbol{\theta})}$$

- ▶ You can efficiently implement the objective function “loss” given your data and your features  $\boldsymbol{\phi}$ .
- ▶ Because it is continuous and differentiable, and the optimization problem is unconstrained, you can use the *gradient* of loss to iteratively move closer to a minimum.
- ▶ Provable: the function is convex, so these methods will converge to a global minimum. More about this in Eisenstein (2019) Appendix B.

## Practical Point: Computing the Gradient

Deriving the gradient of loss with respect to  $\theta$ , denoted  $\nabla_{\theta}\text{loss}$ , is left as an exercise.

Hint: use the chain rule from calculus and work backward through the computation graph on slide 47.

# Stochastic Gradient Descent

Goal: minimize  $\sum_{i=1}^N g_i(\boldsymbol{\theta})$  with respect to  $\boldsymbol{\theta}$ .

Input: initial value  $\boldsymbol{\theta}$ , number of epochs  $T$ , learning rate  $\alpha$

For  $t \in \{1, \dots, T\}$ :

- ▶ Choose a random permutation  $\pi$  of  $\{1, \dots, N\}$ .
- ▶ For  $i \in \{1, \dots, N\}$ :

$$\boldsymbol{\theta} \leftarrow \boldsymbol{w} - \alpha \cdot \nabla_{\boldsymbol{\theta}} g_{\pi(i)}$$

Output:  $\boldsymbol{\theta}$

# Reflection

We can prove that SGD will eventually get very close to a global minimum of a *convex* objective function. What do you think will happen if we apply SGD to a function that is not convex?

# The Main Dish

# Multinomial Logistic Regression

We can generalize LR to an arbitrary label set  $\mathcal{L}$ .

We need:

1. A more powerful definition of feature functions.
2. An update to the probability distribution.



# Input/Output Features

In LR,  $\phi_j : \mathcal{V}^* \rightarrow \mathbb{R}$  (features only see inputs).

In MLR,  $f_j : \mathcal{V}^* \times \mathcal{L} \rightarrow \mathbb{R}$  (features consider potential output value, too).

► (We deliberately use “ $f$ ” instead of “ $\phi$ ” here.)

General template:

$$f_{\ell, \phi}(\mathbf{x}, y) = \phi(\mathbf{x}) \cdot \mathbf{1}\{y = \ell\}$$

E.g., if  $\mathcal{L} = \{\text{sports, politics, health}\}$ , then we have separate features  $f_{\text{sports, vodka}}^{\text{freq.}}(\mathbf{x}, y)$ ,  $f_{\text{politics, vodka}}^{\text{freq.}}(\mathbf{x}, y)$ , and  $f_{\text{health, vodka}}^{\text{freq.}}(\mathbf{x}, y)$ .

# Multinomial Logistic Regression

A multinomial logistic regression model is defined by:


- ▶ A collection of feature functions, denoted  $f_1, \dots, f_d$ , each mapping  $\mathcal{V}^* \times \mathcal{L} \rightarrow \mathbb{R}$ .
  - ▶ **The designer of the system chooses the features.**
- ▶ A coefficient or “weight” for every feature, denoted  $\theta_1, \dots, \theta_d$ , each  $\in \mathbb{R}$ .
  - ▶ **The weights are “parameters” that are chosen automatically by applying a learning algorithm.**

$$\text{score}_{\text{MLR}}(\mathbf{x}, y; \boldsymbol{\theta}) = \sum_{j=1}^d \theta_j f_j(\mathbf{x}, y) = \boldsymbol{\theta}^\top \mathbf{f}(\mathbf{x}, y)$$

$$\text{classify}_{\text{MLR}}(\mathbf{x}) = \arg \max_{y \in \mathcal{L}} \text{score}_{\text{MLR}}(\mathbf{x}, y; \boldsymbol{\theta})$$


# Geometric View of MLR

Suppose we have instance  $x$ ,  $\mathcal{Y} = \{y_1, y_2, y_3, y_4\}$ , and there are only two features,  $f_1$  and  $f_2$ .



## Geometric View of MLR


Suppose we have instance  $x$ ,  $\mathcal{Y} = \{y_1, y_2, y_3, y_4\}$ , and there are only two features,  $f_1$  and  $f_2$ .



$$\theta \cdot \mathbf{f} = \theta_1 f_1 + \theta_2 f_2 = 0$$

## Geometric View of MLR


Suppose we have instance  $x$ ,  $\mathcal{Y} = \{y_1, y_2, y_3, y_4\}$ , and there are only two features,  $f_1$  and  $f_2$ .



$$\text{distance}(\boldsymbol{\theta} \cdot \mathbf{f} = 0, \mathbf{f}_0) = \frac{|\boldsymbol{\theta} \cdot \mathbf{f}_0|}{\|\boldsymbol{\theta}\|_2} \propto |\boldsymbol{\theta} \cdot \mathbf{f}_0|$$

## Geometric View of MLR


Suppose we have instance  $x$ ,  $\mathcal{Y} = \{y_1, y_2, y_3, y_4\}$ , and there are only two features,  $f_1$  and  $f_2$ .



$$\theta \cdot \mathbf{f}(x, y_1) > \theta \cdot \mathbf{f}(x, y_3) > \theta \cdot \mathbf{f}(x, y_4) > 0 \geq \theta \cdot \mathbf{f}(x, y_2)$$

## Geometric View of MLR


Suppose we have instance  $x$ ,  $\mathcal{Y} = \{y_1, y_2, y_3, y_4\}$ , and there are only two features,  $f_1$  and  $f_2$ .



$$\text{score}(x, y_1) > \text{score}(x, y_3) > \text{score}(x, y_4) > \text{score}(x, y_2)$$

# Geometric View of MLR


Suppose we have instance  $x$ ,  $\mathcal{Y} = \{y_1, y_2, y_3, y_4\}$ , and there are only two features,  $f_1$  and  $f_2$ .





## Geometric View of MLR

Suppose we have instance  $x$ ,  $\mathcal{Y} = \{y_1, y_2, y_3, y_4\}$ , and there are only two features,  $f_1$  and  $f_2$ .



$$\text{score}(x, y_3) > \text{score}(x, y_1) > \text{score}(x, y_2) > \text{score}(x, y_4)$$

# Probabilistic View of MLR

Our model defines a *probability* distribution over the labels  $\mathcal{L}$ .

# Probabilistic View of MLR

Our model defines a *probability* distribution over the labels  $\mathcal{L}$ .

First, we need to introduce a new function from vectors to vectors.

$$\begin{aligned}\text{softmax}(\langle t_1, t_2, \dots, t_k \rangle) &= \left\langle \frac{e^{t_1}}{\sum_{j=1}^k e^{t_j}}, \frac{e^{t_2}}{\sum_{j=1}^k e^{t_j}}, \dots, \frac{e^{t_k}}{\sum_{j=1}^k e^{t_j}} \right\rangle \\ &= \frac{\exp \mathbf{t}}{\|\exp \mathbf{t}\|_1}\end{aligned}$$

Note the use of element-wise exponential:

$$\exp(\mathbf{t}) = \langle \exp t_1, \exp t_2, \dots, \exp t_k \rangle.$$

# Probabilistic View of MLR

Our model defines a *probability* distribution over the labels  $\mathcal{L}$ .

$$p_{\text{MLR}}(Y \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \text{softmax} \left( \langle \text{score}_{\text{MLR}}(\mathbf{x}, \ell; \boldsymbol{\theta}) \rangle_{\ell \in \mathcal{L}} \right)$$

# Probabilistic View of MLR

Our model defines a *probability* distribution over the labels  $\mathcal{L}$ .

$$p_{\text{MLR}}(Y \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \text{softmax} \left( \langle \text{score}_{\text{MLR}}(\mathbf{x}, \ell; \boldsymbol{\theta}) \rangle_{\ell \in \mathcal{L}} \right)$$
$$Z(\mathbf{x}; \boldsymbol{\theta}) = \sum_{\ell' \in \mathcal{L}} \exp \text{score}_{\text{MLR}}(\mathbf{x}, \ell'; \boldsymbol{\theta})$$

# Probabilistic View of MLR

Our model defines a *probability* distribution over the labels  $\mathcal{L}$ .

$$p_{\text{MLR}}(Y \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \text{softmax}(\langle \text{score}_{\text{MLR}}(\mathbf{x}, \ell; \boldsymbol{\theta}) \rangle_{\ell \in \mathcal{L}})$$

$$Z(\mathbf{x}; \boldsymbol{\theta}) = \sum_{\ell' \in \mathcal{L}} \exp \text{score}_{\text{MLR}}(\mathbf{x}, \ell'; \boldsymbol{\theta})$$

$$p_{\text{MLR}}(Y = \ell \mid \mathbf{X} = \mathbf{x}; \boldsymbol{\theta}) = \frac{\exp \text{score}_{\text{MLR}}(\mathbf{x}, \ell; \boldsymbol{\theta})}{Z(\mathbf{x}; \boldsymbol{\theta})}$$

# Probabilistic View of MLR


This slide is almost identical to slide 42!

This suggests using the principle of maximum likelihood to estimate  $\theta$ :

$$\begin{aligned}\theta^* &= \arg \max_{\theta \in \mathbb{R}^d} \prod_{i=1}^n p_{\text{MLR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta) \\ &= \arg \max_{\theta \in \mathbb{R}^d} \sum_{i=1}^n \log p_{\text{MLR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta) \\ &= \arg \min_{\theta \in \mathbb{R}^d} \sum_{i=1}^n \underbrace{-\log p_{\text{MLR}}(Y = y_i \mid \mathbf{X} = \mathbf{x}_i; \theta)}_{\text{sometimes called "log loss" or "cross entropy"}}$$

# Reflection: Computation Graph View of MLR

What do you need to change from the LR case?





# Learning for Multinomial Logistic Regression

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta} \in \mathbb{R}^d} \sum_{i=1}^n - \underbrace{\boldsymbol{\theta}^\top \mathbf{f}(\mathbf{x}_i, y_i)}_{\text{"hope"}} + \underbrace{\log \sum_{\ell \in \mathcal{L}} \exp(\boldsymbol{\theta}^\top \mathbf{f}(\mathbf{x}_i, \ell))}_{\text{"fear"}}$$

See slide 31; all points are the same!

# (M)LR Tends to Overfit

If a particular feature  $f_j$  is usually positive, then it always improves the loss to increase  $\theta_j$ .

Regularization: discourage every  $\theta_j$  from getting too large in magnitude.

# Regularization

$$\arg \min_{\boldsymbol{\theta}} \text{loss}(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_p^p$$

where  $\lambda > 0$  is a “hyperparameter” and  $p = 2$  or  $1$ .

## $\ell_1$ Regularization

This case warrants a little more discussion:

$$\min_{\mathbf{w}} \text{loss}(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1$$

Note that:

$$\|\boldsymbol{\theta}\|_1 = \sum_{j=1}^d |\theta_j|$$

- This results in **sparsity** (i.e., many  $\theta_j = 0$ ).

## $\ell_1$ Regularization

This case warrants a little more discussion:

$$\min_{\mathbf{w}} \text{loss}(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1$$

Note that:

$$\|\boldsymbol{\theta}\|_1 = \sum_{j=1}^d |\theta_j|$$

- ▶ This results in **sparsity** (i.e., many  $\theta_j = 0$ ).
  - ▶ Many have argued that this is a good thing (Tibshirani, 1996); it's a kind of feature selection.

## $\ell_1$ Regularization

This case warrants a little more discussion:

$$\min_{\mathbf{w}} \text{loss}(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1$$

Note that:

$$\|\boldsymbol{\theta}\|_1 = \sum_{j=1}^d |\theta_j|$$

- ▶ This results in **sparsity** (i.e., many  $\theta_j = 0$ ).
  - ▶ Many have argued that this is a good thing (Tibshirani, 1996); it's a kind of feature selection.
  - ▶ Do not confuse it with data *sparseness* (a problem to be overcome)!

## $\ell_1$ Regularization

This case warrants a little more discussion:

$$\min_{\mathbf{w}} \text{loss}(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1$$

Note that:

$$\|\boldsymbol{\theta}\|_1 = \sum_{j=1}^d |\theta_j|$$

- ▶ This results in **sparsity** (i.e., many  $\theta_j = 0$ ).
  - ▶ Many have argued that this is a good thing (Tibshirani, 1996); it's a kind of feature selection.
  - ▶ Do not confuse it with data *sparseness* (a problem to be overcome)!
- ▶ This is not differentiable at  $\theta_j = 0$ .

## $\ell_1$ Regularization

This case warrants a little more discussion:

$$\min_{\mathbf{w}} \text{loss}(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1$$

Note that:


$$\|\boldsymbol{\theta}\|_1 = \sum_{j=1}^d |\theta_j|$$

- ▶ This results in **sparsity** (i.e., many  $\theta_j = 0$ ).
  - ▶ Many have argued that this is a good thing (Tibshirani, 1996); it's a kind of feature selection.
  - ▶ Do not confuse it with data *sparseness* (a problem to be overcome)!
- ▶ This is not differentiable at  $\theta_j = 0$ .
- ▶ Optimization: special solutions for batch (e.g., Andrew and Gao, 2007) and stochastic (e.g., Langford et al., 2009) settings.



# Reflection: Computation Graph View of MLR

What do you need to change for regularization?



# MLR Learning

If we had more time, we'd study this problem more carefully!

Here's what you must remember:

- ▶ There is no closed form for the objective function; you must use a numerical optimization algorithm like stochastic gradient descent.
- ▶ MLR is powerful but expensive ( $Z(\mathbf{x}_i; \boldsymbol{\theta})$ ).
- ▶ Regularization is very important; we don't actually do MLE. If you want to be absolutely precise, you're minimizing the regularized log loss.

# Digestif: Connections

Slight changes to the loss function lead to other well-known learning methods.

- ▶ Perceptron: change “fear” to  $\max_{\ell \in \mathcal{L}} \text{score}(\mathbf{x}, \ell; \boldsymbol{\theta})$
- ▶ Linear support vector machine: change “fear” to  $\max_{\ell \in \mathcal{L}} \text{score}(\mathbf{x}, \ell; \boldsymbol{\theta}) + (\text{cost of substituting } \ell \text{ for } y)$

# Digestif: Connections

Slight changes to the loss function lead to other well-known learning methods.

- ▶ Perceptron: change “fear” to  $\max_{\ell \in \mathcal{L}} \text{score}(\mathbf{x}, \ell; \boldsymbol{\theta})$
- ▶ Linear support vector machine: change “fear” to  $\max_{\ell \in \mathcal{L}} \text{score}(\mathbf{x}, \ell; \boldsymbol{\theta}) + (\text{cost of substituting } \ell \text{ for } y)$

The model I presented as “MLR” has gone by other names:

- ▶ Maximum entropy model, because it is provable that  $p_{\text{MLR}}(Y \mid \mathbf{X}; \boldsymbol{\theta}^*)$  is the distribution with the greatest entropy (uncertainty about  $Y$ ) under the constraint that  $\mathbb{E}_p \mathbf{f} = \tilde{\mathbb{E}} \mathbf{f}$ . See Berger et al. (1996).
- ▶ Exponential model, because it is a member of the generalized exponential family.

# On Data

For machine learning methods, the math can be demanding!

This makes it easy to forget the importance of the data and how we represent it (features).

# On Features

Feature engineering is something some people love and others hate.

# On Features

Feature engineering is something some people love and others hate.

There have been many attempts to automate it, either by throwing in a huge number and letting the learner decide (e.g., via sparse regularization), or searching for new, complex features by combining simpler ones, or learning them “from scratch.”

# On Features

Feature engineering is something some people love and others hate.

There have been many attempts to automate it, either by throwing in a huge number and letting the learner decide (e.g., via sparse regularization), or searching for new, complex features by combining simpler ones, or learning them “from scratch.”

Responsible impact: just because you have excluded features that you don't want your model to know about doesn't mean you've excluded all the *correlates* of those features!



# References I

- Galen Andrew and Jianfeng Gao. Scalable training of  $\ell_1$ -regularized log-linear models. In *Proc. of ICML*, 2007.
- Adam Berger, Stephen Della Pietra, and Vincent Della Pietra. A maximum entropy approach to natural language processing. *Computational Linguistics*, 22(1):39–71, 1996.
- Jacob Eisenstein. *Introduction to Natural Language Processing*. MIT Press, 2019.
- John Langford, Lihong Li, and Tong Zhang. Sparse online learning via truncated gradient. In *NeurIPS*, 2009.
- Quinn McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
- Noah A. Smith. *Linguistic Structure Prediction*. Synthesis Lectures on Human Language Technologies. Morgan and Claypool, 2011. URL <http://www.morganclaypool.com/doi/pdf/10.2200/S00361ED1V01Y201105HLT013.pdf>.
- Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society. Series B (Methodological)*, pages 267–288, 1996.

# Extras

# Cross-Validation

Remember that  $\hat{A}$ ,  $\hat{P}$ ,  $\hat{R}$ , and  $\hat{F}_1$  are all *estimates* of the classifier's quality under the true data distribution.

- ▶ Estimates are noisy!

$K$ -fold cross-validation:

- ▶ Partition the training set into  $K$  non-overlapping “folds”  $\mathbf{x}^1, \dots, \mathbf{x}^K$ .
- ▶ For  $i \in \{1, \dots, K\}$ :
  - ▶ Train on  $\mathbf{x}_{1:n} \setminus \mathbf{x}^i$ , using  $\mathbf{x}^i$  as development data.
  - ▶ Estimate quality on the  $i$ th development set:  $\hat{A}^i$
- ▶ Report the average:

$$\hat{A} = \frac{1}{K} \sum_{i=1}^K \hat{A}^i$$

and perhaps also the standard error.

# Statistical Significance

Suppose we have two classifiers,  $\text{classify}_1$  and  $\text{classify}_2$ .

# Statistical Significance

Suppose we have two classifiers,  $\text{classify}_1$  and  $\text{classify}_2$ .

Is  $\text{classify}_1$  better? The “null hypothesis,” denoted  $H_0$ , is that it isn't. But if  $\hat{A}_1 \gg \hat{A}_2$ , we are tempted to believe otherwise.

# Statistical Significance

Suppose we have two classifiers,  $\text{classify}_1$  and  $\text{classify}_2$ .

Is  $\text{classify}_1$  better? The “null hypothesis,” denoted  $H_0$ , is that it isn't. But if  $\hat{A}_1 \gg \hat{A}_2$ , we are tempted to believe otherwise.

How much larger must  $\hat{A}_1$  be than  $\hat{A}_2$  to *reject*  $H_0$ ?

# Statistical Significance

Suppose we have two classifiers,  $\text{classify}_1$  and  $\text{classify}_2$ .

Is  $\text{classify}_1$  better? The “null hypothesis,” denoted  $H_0$ , is that it isn't. But if  $\hat{A}_1 \gg \hat{A}_2$ , we are tempted to believe otherwise.

How much larger must  $\hat{A}_1$  be than  $\hat{A}_2$  to *reject*  $H_0$ ?

Frequentist view: how (im)probable is the observed difference, given  $H_0 = \text{true}$ ?

# Statistical Significance

Suppose we have two classifiers,  $\text{classify}_1$  and  $\text{classify}_2$ .

Is  $\text{classify}_1$  better? The “null hypothesis,” denoted  $H_0$ , is that it isn't. But if  $\hat{A}_1 \gg \hat{A}_2$ , we are tempted to believe otherwise.

How much larger must  $\hat{A}_1$  be than  $\hat{A}_2$  to *reject*  $H_0$ ?

Frequentist view: how (im)probable is the observed difference, given  $H_0 = \text{true}$ ?

Caution: statistical significance is neither necessary nor sufficient for research significance or practical usefulness!



# A Hypothesis Test for Text Classifiers

McNemar (1947)

1. The null hypothesis:  $A_1 = A_2$
2. Pick significance level  $\alpha$ , an “acceptably” high probability of incorrectly rejecting  $H_0$ .
3. Calculate the test statistic,  $k$  (explained in the next slide).
4. Calculate the probability of a *more extreme* value of  $k$ , assuming  $H_0$  is true; this is the  $p$ -value.
5. Reject the null hypothesis if the  $p$ -value is less than  $\alpha$ .

The  $p$ -value is  $p(\text{this observation} \mid H_0 \text{ is true})$ , not the other way around!

## McNemar's Test: Details

Assumptions: independent (test) samples and binary measurements. Count test set error patterns:

	classify <sub>1</sub> is incorrect	classify <sub>1</sub> is correct	
classify <sub>2</sub> is incorrect	$c_{00}$	$c_{10}$	
classify <sub>2</sub> is correct	$c_{01}$	$c_{11}$	$m \cdot \hat{A}_2$
		$m \cdot \hat{A}_1$	

If  $A_1 = A_2$ , then  $c_{01}$  and  $c_{10}$  are each distributed according to  $\text{Binomial}(c_{01} + c_{10}, \frac{1}{2})$ .

test statistic  $k = \min\{c_{01}, c_{10}\}$

$$p\text{-value} = \frac{1}{2^{c_{01}+c_{10}-1}} \sum_{j=0}^k \binom{c_{01} + c_{10}}{j}$$

# Other Tests

Different tests make different assumptions.

Sometimes we calculate an interval that would be “unsurprising” under  $H_0$  and test whether a test statistic falls in that interval (e.g.,  $t$ -test and Wald test).

In many cases, there is no closed form for estimating  $p$ -values, so we use random approximations (e.g., permutation test and paired bootstrap test).

If you do lots of tests, you need to correct for that!

Read lots more in Smith (2011), appendix B.